

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

In [2]:

```
df=pd.read_csv(r'C:\Users\user\Downloads\fiat500_VehicleSelection_Dataset (2).csv')
df
```

Out[2]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price	Unnamed: 10
0	1.0	lounge	51.0	882.0	25000.0	1.0	44.907242	8.611559868	8900	
1	2.0	pop	51.0	1186.0	32500.0	1.0	45.666359	12.24188995	8800	
2	3.0	sport	74.0	4658.0	142228.0	1.0	45.503300	11.41784	4200	
3	4.0	lounge	51.0	2739.0	160000.0	1.0	40.633171	17.63460922	6000	
4	5.0	pop	73.0	3074.0	106880.0	1.0	41.903221	12.49565029	5700	
...	...	...	...	...	...	...	...	...	...	...
1544	NaN	NaN	NaN	NaN	NaN	NaN	NaN	length	5	
1545	NaN	NaN	NaN	NaN	NaN	NaN	NaN	concat	lonprice	
1546	NaN	NaN	NaN	NaN	NaN	NaN	NaN	Null values	NO	
1547	NaN	NaN	NaN	NaN	NaN	NaN	NaN	find	1	
1548	NaN	NaN	NaN	NaN	NaN	NaN	NaN	search	1	

1549 rows × 11 columns

In [3]:

```
df.head(10)
```

Out[3]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price	Unnamed: 10
0	1.0	lounge	51.0	882.0	25000.0	1.0	44.907242	8.611559868	8900	NaN
1	2.0	pop	51.0	1186.0	32500.0	1.0	45.666359	12.24188995	8800	NaN
2	3.0	sport	74.0	4658.0	142228.0	1.0	45.503300	11.41784	4200	NaN
3	4.0	lounge	51.0	2739.0	160000.0	1.0	40.633171	17.63460922	6000	NaN
4	5.0	pop	73.0	3074.0	106880.0	1.0	41.903221	12.49565029	5700	NaN
5	6.0	pop	74.0	3623.0	70225.0	1.0	45.000702	7.68227005	7900	NaN
6	7.0	lounge	51.0	731.0	11600.0	1.0	44.907242	8.611559868	10750	NaN
7	8.0	lounge	51.0	1521.0	49076.0	1.0	41.903221	12.49565029	9190	NaN
8	9.0	sport	73.0	4049.0	76000.0	1.0	45.548000	11.54946995	5600	NaN
9	10.0	sport	51.0	3653.0	89000.0	1.0	45.438301	10.99170017	6000	NaN

In [4]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1549 entries, 0 to 1548
Data columns (total 11 columns):
 #   Column              Non-Null Count  Dtype  
---  -
 0   ID                   1538 non-null   float64
 1   model                1538 non-null   object  
 2   engine_power         1538 non-null   float64
 3   age_in_days          1538 non-null   float64
 4   km                   1538 non-null   float64
 5   previous_owners      1538 non-null   float64
 6   lat                  1538 non-null   float64
 7   lon                  1549 non-null   object  
 8   price                1549 non-null   object  
 9   Unnamed: 9           0 non-null      float64
10   Unnamed: 10          1 non-null      object  
dtypes: float64(7), object(4)
memory usage: 133.2+ KB
```

In [5]:

```
df.describe()
```

Out[5]:

	ID	engine_power	age_in_days	km	previous_owners	lat	Unnamed: 9
count	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	0.0
mean	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.541361	NaN
std	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.133518	NaN
min	1.000000	51.000000	366.000000	1232.000000	1.000000	36.855839	NaN
25%	385.250000	51.000000	670.000000	20006.250000	1.000000	41.802990	NaN
50%	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.394096	NaN
75%	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.467960	NaN
max	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.795612	NaN

In [6]:

```
df.columns
```

Out[6]:


```
Index(['ID', 'model', 'engine_power', 'age_in_days', 'km', 'previous_owners',
       'lat', 'lon', 'price', 'Unnamed: 9', 'Unnamed: 10'],
      dtype='object')
```

In [7]:

```
sns.pairplot(df)
```

Out[7]:

<seaborn.axisgrid.PairGrid at 0x2b1837ceee0>



In [8]:


```
sns.distplot(df['km'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

Out[8]:

```
<AxesSubplot:xlabel='km', ylabel='Density'>
```




In [9]:

```
sns.displot(df["km"])
```

Out[9]:

```
<seaborn.axisgrid.FacetGrid at 0x2b18556bb20>
```



In [10]:


```
df1=df[['ID', 'model', 'engine_power', 'age_in_days', 'km', 'previous_owners',  
        'lat', 'lon', 'price', 'Unnamed: 9', 'Unnamed: 10']]
```

In [11]:

```
sns.heatmap(df1.corr())
```

Out[11]:

<AxesSubplot:>



In [12]:

```
x=df1[['ID', 'engine_power', 'age_in_days', 'previous_owners', 'lat', 'Unnamed: 9']]  
y=df1[['km']]
```

In [13]:

```
from sklearn.model_selection import train_test_split
```

In [14]:

```
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```

In [20]:

```
from sklearn.linear_model import LinearRegression
```

```
lr=LinearRegression()
```

```
lr.fit(x_train,y_train)#ValueError: Input contains NaN, infinity or a value too large for dtype('float64')
```

```

-----
ValueError                                Traceback (most recent call last)
<ipython-input-20-58bc3c346312> in <module>
      2
      3 lr=LinearRegression()
----> 4 lr.fit(x_train,y_train)#ValueError: Input contains NaN, infinity or a value too large for dtype('float64')

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\base.py in fit(self, X, y, sample_weight)
    516         accept_sparse = False if self.positive else ['csr', 'csc', 'coo']
    517
--> 518         X, y = self._validate_data(X, y, accept_sparse=accept_sparse,
    519                                     y_numeric=True, multi_output=True)
    520

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\base.py in _validate_data(self, X, y, reset, validate_separately, **check_params)
    431         y = check_array(y, **check_y_params)
    432     else:
--> 433         X, y = check_X_y(X, y, **check_params)
    434         out = X, y
    435

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in inner_f(*args, **kwargs)
    61         extra_args = len(args) - len(all_args)
    62         if extra_args <= 0:
----> 63             return f(*args, **kwargs)
    64
    65         # extra_args > 0

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in check_X_y(X, y, accept_sparse, accept_large_sparse, dtype, order, copy, force_all_finite, ensure_2d, allow_nd, multi_output, ensure_min_samples, ensure_min_features, y_numeric, estimator)
    812         raise ValueError("y cannot be None")
    813
--> 814         X = check_array(X, accept_sparse=accept_sparse,
    815                         accept_large_sparse=accept_large_sparse,
    816                         dtype=dtype, order=order, copy=copy,

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in inner_f(*args, **kwargs)
    61         extra_args = len(args) - len(all_args)
    62         if extra_args <= 0:
----> 63             return f(*args, **kwargs)
    64
    65         # extra_args > 0

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in check_array(array, accept_sparse, accept_large_sparse, dtype, order, copy, force_all_finite, ensure_2d, allow_nd, ensure_min_samples, ensure_min_features, estimator)
    661
    662         if force_all_finite:
--> 663             _assert_all_finite(array,
    664                               allow_nan=force_all_finite == 'allow-nan')
    665

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in _assert_all_finite(X, allow_nan, msg_dtype)
    101         not allow_nan and not np.isfinite(X).all()):
    102         type_err = 'infinity' if allow_nan else 'NaN, infinity'
--> 103         raise ValueError(
    104             msg_err.format
    105             (type_err,

ValueError: Input contains NaN, infinity or a value too large for dtype('float64').

```

In [16]:

```
print(lr.intercept_)
```

```
-----  
AttributeError                                Traceback (most recent call last)  
<ipython-input-16-182bb45ab960> in <module>  
----> 1 print(lr.intercept_)
```

**AttributeError:** 'LinearRegression' object has no attribute 'intercept\_'

In [17]:

```
coef= pd.DataFrame(lr.coef_)  
coef
```

```
-----  
AttributeError                                Traceback (most recent call last)  
<ipython-input-17-0ff321f0a2a5> in <module>  
----> 1 coef= pd.DataFrame(lr.coef_)  
      2 coef
```

**AttributeError:** 'LinearRegression' object has no attribute 'coef\_'



In [18]:

```
print(lr.score(x_test,y_test))
```

```
-----
NotFittedError                                Traceback (most recent call last)
<ipython-input-18-6bc23016a4ce> in <module>
----> 1 print(lr.score(x_test,y_test))

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\base.py in score(self, X, y, sample_weight)
    551
    552         from .metrics import r2_score
--> 553         y_pred = self.predict(X)
    554         return r2_score(y, y_pred, sample_weight=sample_weight)
    555

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\base.py in predict(self, X)
    236         Returns predicted values.
    237         """
--> 238         return self._decision_function(X)
    239
    240     _preprocess_data = staticmethod(_preprocess_data)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\base.py in _decision_function(self, X)
    216
    217     def _decision_function(self, X):
--> 218         check_is_fitted(self)
    219
    220         X = check_array(X, accept_sparse=['csr', 'csc', 'coo'])

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in inner_f(*args, **kwargs)
    61         extra_args = len(args) - len(all_args)
    62         if extra_args <= 0:
--> 63             return f(*args, **kwargs)
    64
    65         # extra_args > 0

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in check_is_fitted(estimator, attributes, msg, all_or_any)
   1039
   1040     if not attrs:
-> 1041         raise NotFittedError(msg % {'name': type(estimator).__name__})
   1042
   1043

NotFittedError: This LinearRegression instance is not fitted yet. Call 'fit' with appropriate arguments before using this estimator.
```

In [19]:

```
prediction = lr.predict(x_test)
plt.scatter(y_test, prediction)
```

```
-----
NotFittedError                                Traceback (most recent call last)
<ipython-input-19-10d398fd7dc3> in <module>
----> 1 prediction = lr.predict(x_test)
      2 plt.scatter(y_test, prediction)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\base.py in predict(self,
X)
    236         Returns predicted values.
    237         """
--> 238         return self._decision_function(X)
    239
    240     _preprocess_data = staticmethod(_preprocess_data)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\base.py in _decision_funct
ion(self, X)
    216
    217     def _decision_function(self, X):
--> 218         check_is_fitted(self)
    219
    220         X = check_array(X, accept_sparse=['csr', 'csc', 'coo'])

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in inner_f(*args, **
kwargs)
    61         extra_args = len(args) - len(all_args)
    62         if extra_args <= 0:
--> 63             return f(*args, **kwargs)
    64
    65         # extra_args > 0

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in check_is_fitted(e
stimator, attributes, msg, all_or_any)
    1039
    1040     if not attrs:
-> 1041         raise NotFittedError(msg % {'name': type(estimator).__name__})
    1042
    1043

NotFittedError: This LinearRegression instance is not fitted yet. Call 'fit' with appropria
te arguments before using this estimator.
```