

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```


In [2]:

```
df=pd.read_csv(r'C:\Users\user\Downloads\6_Salesworkload1.csv')
df
```

Out[2]:

	MonthYear	Time index	Country	StoreID	City	Dept_ID	Dept. Name	HoursOwn	Hour
0	10.2016	1.0	United Kingdom	88253.0	London (I)	1.0	Dry	3184.764	
1	10.2016	1.0	United Kingdom	88253.0	London (I)	2.0	Frozen	1582.941	
2	10.2016	1.0	United Kingdom	88253.0	London (I)	3.0	other	47.205	
3	10.2016	1.0	United Kingdom	88253.0	London (I)	4.0	Fish	1623.852	
4	10.2016	1.0	United Kingdom	88253.0	London (I)	5.0	Fruits & Vegetables	1759.173	
...	
7653	06.2017	9.0	Sweden	29650.0	Gothenburg	12.0	Checkout	6322.323	
7654	06.2017	9.0	Sweden	29650.0	Gothenburg	16.0	Customer Services	4270.479	
7655	06.2017	9.0	Sweden	29650.0	Gothenburg	11.0	Delivery	0	
7656	06.2017	9.0	Sweden	29650.0	Gothenburg	17.0	others	2224.929	
7657	06.2017	9.0	Sweden	29650.0	Gothenburg	18.0	all	39652.2	

7658 rows × 14 columns



In [3]:

```
df.head(10)
```

Out[3]:

	MonthYear	Time index	Country	StoreID	City	Dept_ID	Dept. Name	HoursOwn	HoursLease
0	10.2016	1.0	United Kingdom	88253.0	London (I)	1.0	Dry	3184.764	0.0
1	10.2016	1.0	United Kingdom	88253.0	London (I)	2.0	Frozen	1582.941	0.0
2	10.2016	1.0	United Kingdom	88253.0	London (I)	3.0	other	47.205	0.0
3	10.2016	1.0	United Kingdom	88253.0	London (I)	4.0	Fish	1623.852	0.0
4	10.2016	1.0	United Kingdom	88253.0	London (I)	5.0	Fruits & Vegetables	1759.173	0.0
5	10.2016	1.0	United Kingdom	88253.0	London (I)	6.0	Meat	8270.316	0.0
6	10.2016	1.0	United Kingdom	88253.0	London (I)	13.0	Food	16468.251	0.0
7	10.2016	1.0	United Kingdom	88253.0	London (I)	7.0	Clothing	4698.471	0.0
8	10.2016	1.0	United Kingdom	88253.0	London (I)	8.0	Household	1183.272	0.0
9	10.2016	1.0	United Kingdom	88253.0	London (I)	9.0	Hardware	2029.815	0.0

In [4]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7658 entries, 0 to 7657
Data columns (total 14 columns):
#   Column          Non-Null Count  Dtype
---  -
0   MonthYear       7658 non-null   object
1   Time index      7650 non-null   float64
2   Country         7650 non-null   object
3   StoreID         7650 non-null   float64
4   City            7650 non-null   object
5   Dept_ID         7650 non-null   float64
6   Dept. Name      7650 non-null   object
7   HoursOwn        7650 non-null   object
8   HoursLease      7650 non-null   float64
9   Sales units     7650 non-null   float64
10  Turnover        7650 non-null   float64
11  Customer         0 non-null      float64
12  Area (m2)       7650 non-null   object
13  Opening hours   7650 non-null   object
dtypes: float64(7), object(7)
memory usage: 837.7+ KB
```

In [5]:

```
df.describe()
```

Out[5]:

	Time index	StoreID	Dept_ID	HoursLease	Sales units	Turnover	Cu
count	7650.000000	7650.000000	7650.000000	7650.000000	7.650000e+03	7.650000e+03	
mean	5.000000	61995.220000	9.470588	22.036078	1.076471e+06	3.721393e+06	
std	2.582158	29924.581631	5.337429	133.299513	1.728113e+06	6.003380e+06	
min	1.000000	12227.000000	1.000000	0.000000	0.000000e+00	0.000000e+00	
25%	3.000000	29650.000000	5.000000	0.000000	5.457125e+04	2.726798e+05	
50%	5.000000	75400.500000	9.000000	0.000000	2.932300e+05	9.319575e+05	
75%	7.000000	87703.000000	14.000000	0.000000	9.175075e+05	3.264432e+06	
max	9.000000	98422.000000	18.000000	3984.000000	1.124296e+07	4.271739e+07	

In [6]:

```
df.columns
```

Out[6]:

```
Index(['MonthYear', 'Time index', 'Country', 'StoreID', 'City', 'Dept_ID',  
      'Dept. Name', 'HoursOwn', 'HoursLease', 'Sales units', 'Turnover',  
      'Customer', 'Area (m2)', 'Opening hours'],  
      dtype='object')
```

In [7]:

```
sns.pairplot(df)
```

Out[7]:

<seaborn.axisgrid.PairGrid at 0x1f5b1ca83a0>



In [8]:


```
sns.distplot(df['HoursLease'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557:
FutureWarning: `distplot` is a deprecated function and will be removed in
a future version. Please adapt your code to use either `displot` (a figure
-level function with similar flexibility) or `histplot` (an axes-level fun
ction for histograms).

```
warnings.warn(msg, FutureWarning)
```

Out[8]:

<AxesSubplot:xlabel='HoursLease', ylabel='Density'>




In [9]:

```
sns.displot(df["HoursLease"])
```

Out[9]:

<seaborn.axisgrid.FacetGrid at 0x1f5acef15e0>



In [10]:


```
df1=df[['MonthYear', 'Time index', 'Country', 'StoreID', 'City', 'Dept_ID',
        'Dept. Name', 'HoursOwn', 'HoursLease', 'Sales units', 'Turnover',
        'Customer', 'Area (m2)', 'Opening hours']]
```

In [11]:

```
sns.heatmap(df1.corr())
```

Out[11]:

<AxesSubplot:>



In [16]:

```
x=df1[['Time index', 'StoreID', 'Dept_ID',
        'HoursLease', 'Sales units', 'Turnover',
        'Customer']]
y=df1[['HoursLease']]
```

In [17]:

```
from sklearn.model_selection import train_test_split
```

In [18]:

```
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```

In [20]:

```
from sklearn.linear_model import LinearRegression  
  
lr=LinearRegression()  
lr.fit(x_train,y_train)#ValueError: Input contains NaN, infinity or a value too large for
```

```

-----
-
ValueError                                Traceback (most recent call las
t)
<ipython-input-20-58bc3c346312> in <module>
      2
      3 lr=LinearRegression()
----> 4 lr.fit(x_train,y_train)#ValueError: Input contains NaN, infinity o
r a value too large for dtype('float64')

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\_base.py i
n fit(self, X, y, sample_weight)
    516         accept_sparse = False if self.positive else ['csr', 'csc',
'coo']
    517
--> 518         X, y = self._validate_data(X, y, accept_sparse=accept_spar
se,
    519                                     y_numeric=True, multi_output=Tr
ue)
    520

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\base.py in _validate_da
ta(self, X, y, reset, validate_separately, **check_params)
    431         y = check_array(y, **check_y_params)
    432     else:
--> 433         X, y = check_X_y(X, y, **check_params)
    434         out = X, y
    435

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in
inner_f(*args, **kwargs)
    61         extra_args = len(args) - len(all_args)
    62         if extra_args <= 0:
--> 63             return f(*args, **kwargs)
    64
    65         # extra_args > 0

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in
check_X_y(X, y, accept_sparse, accept_large_sparse, dtype, order, copy, fo
rce_all_finite, ensure_2d, allow_nd, multi_output, ensure_min_samples, ens
ure_min_features, y_numeric, estimator)
    812         raise ValueError("y cannot be None")
    813
--> 814         X = check_array(X, accept_sparse=accept_sparse,
    815                           accept_large_sparse=accept_large_sparse,
    816                           dtype=dtype, order=order, copy=copy,

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in
inner_f(*args, **kwargs)
    61         extra_args = len(args) - len(all_args)
    62         if extra_args <= 0:
--> 63             return f(*args, **kwargs)
    64
    65         # extra_args > 0

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in
check_array(array, accept_sparse, accept_large_sparse, dtype, order, copy,
force_all_finite, ensure_2d, allow_nd, ensure_min_samples, ensure_min_feat
ures, estimator)
    661
    662         if force_all_finite:

```



```
--> 663         _assert_all_finite(array,
664                             allow_nan=force_all_finite == 'allo
w-nan')
665
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in
_assert_all_finite(X, allow_nan, msg_dtype)
101         not allow_nan and not np.isfinite(X).all()):
102         type_err = 'infinity' if allow_nan else 'NaN, infinit
y'
--> 103         raise ValueError(
104             msg_err.format
105             (type_err,
```

ValueError: Input contains NaN, infinity or a value too large for dtype('float64').

In [21]:

```
print(lr.intercept_)
```

```
-----
-
AttributeError                                Traceback (most recent call las
t)
<ipython-input-21-182bb45ab960> in <module>
----> 1 print(lr.intercept_)
```

AttributeError: 'LinearRegression' object has no attribute 'intercept_'

In [22]:

```
coef= pd.DataFrame(lr.coef_)
coef
```

```
-----
-
AttributeError                                Traceback (most recent call las
t)
<ipython-input-22-0ff321f0a2a5> in <module>
----> 1 coef= pd.DataFrame(lr.coef_)
      2 coef
```

AttributeError: 'LinearRegression' object has no attribute 'coef_'

In [23]:

```
print(lr.score(x_test,y_test))
```

```
-----
-
NotFittedError                                Traceback (most recent call last)
<ipython-input-23-6bc23016a4ce> in <module>
----> 1 print(lr.score(x_test,y_test))

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\base.py in score(self,
X, y, sample_weight)
    551
    552         from .metrics import r2_score
--> 553         y_pred = self.predict(X)
    554         return r2_score(y, y_pred, sample_weight=sample_weight)
    555

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_base.py in
predict(self, X)
    236         Returns predicted values.
    237         """
--> 238         return self._decision_function(X)
    239
    240         _preprocess_data = staticmethod(_preprocess_data)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_base.py in
_decision_function(self, X)
    216
    217     def _decision_function(self, X):
--> 218         check_is_fitted(self)
    219
    220         X = check_array(X, accept_sparse=['csr', 'csc', 'coo'])

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in
inner_f(*args, **kwargs)
    61         extra_args = len(args) - len(all_args)
    62         if extra_args <= 0:
--> 63             return f(*args, **kwargs)
    64
    65             # extra_args > 0

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in
check_is_fitted(estimator, attributes, msg, all_or_any)
   1039
   1040     if not attrs:
-> 1041         raise NotFittedError(msg % {'name': type(estimator).__name
__})
   1042
   1043

NotFittedError: This LinearRegression instance is not fitted yet. Call 'fit'
with appropriate arguments before using this estimator.
```

In [24]:

```
prediction = lr.predict(x_test)
plt.scatter(y_test, prediction)
```

```
-----
-
NotFittedError                                Traceback (most recent call las
t)
<ipython-input-24-10d398fd7dc3> in <module>
----> 1 prediction = lr.predict(x_test)
      2 plt.scatter(y_test, prediction)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\_base.py i
n predict(self, X)
    236         Returns predicted values.
    237         """
--> 238         return self._decision_function(X)
    239
    240         _preprocess_data = staticmethod(_preprocess_data)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\_base.py i
n _decision_function(self, X)
    216
    217     def _decision_function(self, X):
--> 218         check_is_fitted(self)
    219
    220         X = check_array(X, accept_sparse=['csr', 'csc', 'coo'])

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in
inner_f(*args, **kwargs)
    61         extra_args = len(args) - len(all_args)
    62         if extra_args <= 0:
--> 63             return f(*args, **kwargs)
    64
    65             # extra_args > 0

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in
check_is_fitted(estimator, attributes, msg, all_or_any)
    1039
    1040     if not attrs:
-> 1041         raise NotFittedError(msg % {'name': type(estimator).__name
__})
    1042
    1043

NotFittedError: This LinearRegression instance is not fitted yet. Call 'fi
t' with appropriate arguments before using this estimator.
```